

Qualifying Exam

November 7, 2025

Sam Buxbaum

Motivation

Motivation

Secure computation has taken great leaps in recent years...

Motivation

Secure computation has taken great leaps in recent years...

but it's still impractical in many scenarios.

Motivation

Secure computation has taken great leaps in recent years...

but it's still impractical in many scenarios.

Goal: Fill in the gaps that prevent the use of MPC in practice

Scaling Up MPC

Scaling Up MPC

1. More sophisticated functions

Scaling Up MPC

1. More sophisticated functions

- Oblivious shuffling

Scaling Up MPC

1. More sophisticated functions

- Oblivious shuffling
- Improving the expressivity of MPC

Scaling Up MPC

1. More sophisticated functions

- Oblivious shuffling
- Improving the expressivity of MPC

2. Directly incorporating more parties

Scaling Up MPC

1. More sophisticated functions

- Oblivious shuffling
- Improving the expressivity of MPC

2. Directly incorporating more parties

- Significant practical challenges

Scaling Up MPC

1. More sophisticated functions

- Oblivious shuffling
- Improving the expressivity of MPC

2. Directly incorporating more parties

- Significant practical challenges
- Distribute trust more widely

Oblivious Shuffle Preprocessing

Oblivious Shuffling

Two parties, P_0 and P_1 , want to shuffle a secret-shared vector $[\mathbf{x}]$ by a permutation π .

Oblivious Shuffling

Two parties, P_0 and P_1 , want to shuffle a secret-shared vector $[\mathbf{x}]$ by a permutation π .

- Neither party should learn anything about π

Oblivious Shuffling


Two parties, P_0 and P_1 , want to shuffle a secret-shared vector $[\mathbf{x}]$ by a permutation π .

- Neither party should learn anything about π
- The output is a secret sharing $[\pi(\mathbf{x})]$

Example: Oblivious Sorting

Example: Oblivious Sorting


Oblivious sorting networks require $O(n \log^2 n)$ comparison gates.



Example: Oblivious Sorting

Oblivious sorting networks require $O(n \log^2 n)$ comparison gates.

- We want to match the plaintext $O(n \log n)$ algorithms




Example: Oblivious Sorting

Oblivious sorting networks require $O(n \log^2 n)$ comparison gates.

- We want to match the plaintext $O(n \log n)$ algorithms


The shuffle-then-sort paradigm (Hamada et al.)



Example: Oblivious Sorting

Oblivious sorting networks require $O(n \log^2 n)$ comparison gates.

- We want to match the plaintext $O(n \log n)$ algorithms




The shuffle-then-sort paradigm (Hamada et al.)

- Obviously shuffle, then use a data-dependent sorting algorithm

Example: Oblivious Sorting

Oblivious sorting networks require $O(n \log^2 n)$ comparison gates.

- We want to match the plaintext $O(n \log n)$ algorithms




The shuffle-then-sort paradigm (Hamada et al.)

- Obviously shuffle, then use a data-dependent sorting algorithm
- The revealed permutation is uncorrelated with the input

Example: Oblivious Sorting

Oblivious sorting networks require $O(n \log^2 n)$ comparison gates.


- We want to match the plaintext $O(n \log n)$ algorithms



The shuffle-then-sort paradigm (Hamada et al.)


- Obviously shuffle, then use a data-dependent sorting algorithm
- The revealed permutation is uncorrelated with the input
- Achieves $O(n \log n)$ secure comparisons

Example: Graph Algorithms



Example: Graph Algorithms


Not easy to efficiently express a graph obliviously.



Example: Graph Algorithms

Not easy to efficiently express a graph obliviously.

- Lots of data-dependent operations




Example: Graph Algorithms

Not easy to efficiently express a graph obliviously.

- Lots of data-dependent operations

The GraphSC paradigm (Nayak et al.)




Example: Graph Algorithms

Not easy to efficiently express a graph obliviously.

- Lots of data-dependent operations

The GraphSC paradigm (Nayak et al.)

- Encode the graph as a single list




Example: Graph Algorithms

Not easy to efficiently express a graph obliviously.

- Lots of data-dependent operations

The GraphSC paradigm (Nayak et al.)

- Encode the graph as a single list
- The list has two orderings




Example: Graph Algorithms

Not easy to efficiently express a graph obliviously.

- Lots of data-dependent operations

The GraphSC paradigm (Nayak et al.)

- Encode the graph as a single list
- The list has two orderings
- Two important operations, scatter and gather




Example: Graph Algorithms

Not easy to efficiently express a graph obliviously.

- Lots of data-dependent operations

The GraphSC paradigm (Nayak et al.)

- Encode the graph as a single list
- The list has two orderings
- Two important operations, scatter and gather
- Each operation is efficiently only in one of the orderings




Example: Graph Algorithms

Not easy to efficiently express a graph obliviously.

- Lots of data-dependent operations

The GraphSC paradigm (Nayak et al.)

- Encode the graph as a single list
- The list has two orderings
- Two important operations, scatter and gather
- Each operation is efficiently only in one of the orderings
- Shuffling allows us to convert between the orderings



Broader Applications

Broader Applications

Combine aspects of the circuit model and RAM model of computation.

Broader Applications

Combine aspects of the circuit model and RAM model of computation.

- A bridge between the two worlds

Broader Applications

Combine aspects of the circuit model and RAM model of computation.

- A bridge between the two worlds
- Expressivity makes it central to complex workloads

Broader Applications

Combine aspects of the circuit model and RAM model of computation.

- A bridge between the two worlds
- Expressivity makes it central to complex workloads

How do we do it?

Break the Problem in Half

Break the Problem in Half

- Let $\pi = \pi_1 \circ \pi_0$ be a composition of random permutations

Break the Problem in Half

- Let $\pi = \pi_1 \circ \pi_0$ be a composition of random permutations
 - P_0 knows π_0

Break the Problem in Half

- Let $\pi = \pi_1 \circ \pi_0$ be a composition of random permutations
 - P_0 knows π_0
 - P_1 knows π_1

Break the Problem in Half


- Let $\pi = \pi_1 \circ \pi_0$ be a composition of random permutations
 - P_0 knows π_0
 - P_1 knows π_1
 - Neither party knows anything about π

Break the Problem in Half

- Let $\pi = \pi_1 \circ \pi_0$ be a composition of random permutations
 - P_0 knows π_0
 - P_1 knows π_1
 - Neither party knows anything about π
- One party acts as a sender, the other as a receiver


Break the Problem in Half

- Let $\pi = \pi_1 \circ \pi_0$ be a composition of random permutations
 - P_0 knows π_0
 - P_1 knows π_1
 - Neither party knows anything about π
- One party acts as a sender, the other as a receiver



Break the Problem in Half

- Let $\pi = \pi_1 \circ \pi_0$ be a composition of random permutations
 - P_0 knows π_0
 - P_1 knows π_1
 - Neither party knows anything about π
- One party acts as a sender, the other as a receiver
- Known as "Permute+Share" (Chase et al.)



Preprocessing

Preprocessing

Permute+Share can be broken into preprocessing and online phases.

Preprocessing

Permute+Share can be broken into preprocessing and online phases.

- The preprocessing phase shuffles pseudorandom values

Preprocessing

Permute+Share can be broken into preprocessing and online phases.

- The preprocessing phase shuffles pseudorandom values
- The online phase derandomizes in a single message

Preprocessing

Permute+Share can be broken into preprocessing and online phases.

- The preprocessing phase shuffles pseudorandom values
- The online phase derandomizes in a single message

Preprocessing produces a *permutation correlation*.

Preprocessing

Permute+Share can be broken into preprocessing and online phases.

- The preprocessing phase shuffles pseudorandom values
- The online phase derandomizes in a single message

Preprocessing produces a *permutation correlation*.

$$(\pi, C), (A, B)$$

Preprocessing

Permute+Share can be broken into preprocessing and online phases.

- The preprocessing phase shuffles pseudorandom values
- The online phase derandomizes in a single message

Preprocessing produces a *permutation correlation*.

$$(\pi, C), (A, B)$$

$$\pi(A) = B + C$$

Online Derandomization Protocol

Online Derandomization Protocol

S: (π, C)

R: $(\mathbf{x}, (A, B))$

$\mathbf{x} - A$



$\pi(\mathbf{x} - A) + C$

B

Online Derandomization Protocol

S: (π, C)

R: $(\mathbf{x}, (A, B))$

$\xleftarrow{\mathbf{x} - A}$

$\pi(\mathbf{x} - A) + C$

B

$$\begin{aligned}\text{output} &= \pi(\mathbf{x} - A) + C + B \\ &= \pi(\mathbf{x} - A) + \pi(A) \\ &= \pi(\mathbf{x}) - \pi(A) + \pi(A) \\ &= \pi(\mathbf{x})\end{aligned}$$

Formally Defining the Problem

Formally Defining the Problem

Functionality $\mathcal{F}_{\text{GenPerm}}$

PUBLIC PARAMETERS: Number of elements n and bitwidth ℓ .

INPUT: None

OUTPUT: The functionality samples a uniformly random permutation $\pi \leftarrow S_n$ and uniformly random vectors $A, B, C \in \{0, 1\}^{n \times \ell}$, such that $\pi(A) = B + C$. $\mathcal{F}_{\text{GenPerm}}$ outputs (π, C) to the sender and (A, B) to the receiver.

Formally Defining the Problem

Functionality $\mathcal{F}_{\text{GenPerm}}$

PUBLIC PARAMETERS: Number of elements n and bitwidth ℓ .

INPUT: None

OUTPUT: The functionality samples a uniformly random permutation $\pi \leftarrow S_n$ and uniformly random vectors $A, B, C \in \{0, 1\}^{n \times \ell}$, such that $\pi(A) = B + C$. $\mathcal{F}_{\text{GenPerm}}$ outputs (π, C) to the sender and (A, B) to the receiver.

Three parameters we care about:

Formally Defining the Problem

Functionality $\mathcal{F}_{\text{GenPerm}}$

PUBLIC PARAMETERS: Number of elements n and bitwidth ℓ .

INPUT: None

OUTPUT: The functionality samples a uniformly random permutation $\pi \leftarrow S_n$ and uniformly random vectors $A, B, C \in \{0, 1\}^{n \times \ell}$, such that $\pi(A) = B + C$. $\mathcal{F}_{\text{GenPerm}}$ outputs (π, C) to the sender and (A, B) to the receiver.

Three parameters we care about:

- n , the number of elements in a permutation correlation

Formally Defining the Problem

Functionality $\mathcal{F}_{\text{GenPerm}}$

PUBLIC PARAMETERS: Number of elements n and bitwidth ℓ .

INPUT: None

OUTPUT: The functionality samples a uniformly random permutation $\pi \leftarrow S_n$ and uniformly random vectors $A, B, C \in \{0, 1\}^{n \times \ell}$, such that $\pi(A) = B + C$. $\mathcal{F}_{\text{GenPerm}}$ outputs (π, C) to the sender and (A, B) to the receiver.

Three parameters we care about:

- n , the number of elements in a permutation correlation
- ℓ , the bitwidth of each element of A , B , and C in a single permutation correlation

Formally Defining the Problem

Functionality $\mathcal{F}_{\text{GenPerm}}$

PUBLIC PARAMETERS: Number of elements n and bitwidth ℓ .

INPUT: None

OUTPUT: The functionality samples a uniformly random permutation $\pi \leftarrow S_n$ and uniformly random vectors $A, B, C \in \{0, 1\}^{n \times \ell}$, such that $\pi(A) = B + C$. $\mathcal{F}_{\text{GenPerm}}$ outputs (π, C) to the sender and (A, B) to the receiver.

Three parameters we care about:

- n , the number of elements in a permutation correlation
- ℓ , the bitwidth of each element of A , B , and C in a single permutation correlation
- p , the total number of permutation correlations generated

State of the Art

State of the Art

Peceny, Raghuraman, Rindal, and Shah (PKC 2025)

Efficient Permutation Correlations and Batched Random Access for Two-Party Computation

Stanislav Peceny^{1*}, Srinivasan Raghuraman², Peter Rindal³, and Harshal Shah³

¹ Georgia Institute of Technology

² Visa Research and MIT

³ Visa Research

Abstract. In this work we formalize the notion of a two-party permutation correlation $(A, B), (C, \pi)$ s.t. $\pi(A) = B + C$ for a random permutation π of n elements and vectors $A, B, C \in \mathbb{F}^n$. This correlation can be viewed as an abstraction and generalization of the Chase et al. (Asiacrypt 2020) share translation protocol. We give a systematization of knowledge for how such a permutation correlation can be derandomized

State of the Art

Peceny, Raghuraman, Rindal, and Shah (PKC 2025)

- Two constructions

Efficient Permutation Correlations and Batched Random Access for Two-Party Computation

Stanislav Peceny^{1*}, Srinivasan Raghuraman², Peter Rindal³, and Harshal Shah³

¹ Georgia Institute of Technology

² Visa Research and MIT

³ Visa Research

Abstract. In this work we formalize the notion of a two-party permutation correlation $(A, B), (C, \pi)$ s.t. $\pi(A) = B + C$ for a random permutation π of n elements and vectors $A, B, C \in \mathbb{F}^n$. This correlation can be viewed as an abstraction and generalization of the Chase et al. (Asiacrypt 2020) share translation protocol. We give a systematization of knowledge for how such a permutation correlation can be derandomized

State of the Art

Peceny, Raghuraman, Rindal, and Shah (PKC 2025)

- Two constructions
- Both based on oblivious pseudorandom functions

Efficient Permutation Correlations and Batched Random Access for Two-Party Computation

Stanislav Peceny^{1*}, Srinivasan Raghuraman², Peter Rindal³, and Harshal Shah³

¹ Georgia Institute of Technology

² Visa Research and MIT

³ Visa Research

Abstract. In this work we formalize the notion of a two-party permutation correlation $(A, B), (C, \pi)$ s.t. $\pi(A) = B + C$ for a random permutation π of n elements and vectors $A, B, C \in \mathbb{F}^n$. This correlation can be viewed as an abstraction and generalization of the Chase et al. (Asiacrypt 2020) share translation protocol. We give a systematization of knowledge for how such a permutation correlation can be derandomized

State of the Art

Peceny, Raghuraman, Rindal, and Shah (PKC 2025)

- Two constructions
- Both based on oblivious pseudorandom functions

Construction 1	Construction 2
$O(n\ell)$	$O(n \log \ell)$

Efficient Permutation Correlations and Batched Random Access for Two-Party Computation

Stanislav Peceny^{1*}, Srinivasan Raghuraman², Peter Rindal³, and Harshal Shah³


¹ Georgia Institute of Technology

² Visa Research and MIT

³ Visa Research


Abstract. In this work we formalize the notion of a two-party permutation correlation $(A, B), (C, \pi)$ s.t. $\pi(A) = B + C$ for a random permutation π of n elements and vectors $A, B, C \in \mathbb{F}^n$. This correlation can be viewed as an abstraction and generalization of the Chase et al. (Asiacrypt 2020) share translation protocol. We give a systematization of knowledge for how such a permutation correlation can be derandomized

Oblivious Pseudorandom Functions (OPRFs)



Oblivious Pseudorandom Functions (OPRFs)


A core building block of each construction is an OPRF with secret-shared output.



Oblivious Pseudorandom Functions (OPRFs)

A core building block of each construction is an OPRF with secret-shared output.


- Sender provides an input x



Oblivious Pseudorandom Functions (OPRFs)

A core building block of each construction is an OPRF with secret-shared output.


- Sender provides an input x
- Receiver provides a key k



Oblivious Pseudorandom Functions (OPRFs)

A core building block of each construction is an OPRF with secret-shared output.


- Sender provides an input x
- Receiver provides a key k
- Instantiated with the alternating-moduli family of PRFs




Oblivious Pseudorandom Functions (OPRFs)

A core building block of each construction is an OPRF with secret-shared output.

- Sender provides an input x
- Receiver provides a key k
- Instantiated with the alternating-moduli family of PRFs
- Outputs a secret sharing




Primary Construction




Primary Construction

- Parties agree on a random vector $\mathbf{x}$ by computing $x_i = H(i)$




Primary Construction

- Parties agree on a random vector $\mathbf{x}$ by computing $x_i = H(i)$
- S samples $\pi \leftarrow S_n$




Primary Construction

- Parties agree on a random vector $\mathbf{x}$ by computing $x_i = H(i)$
- S samples $\pi \leftarrow S_n$
- S permutes $\mathbf{x}$ by π to set OPRF inputs




Primary Construction

- Parties agree on a random vector $\mathbf{x}$ by computing $x_i = H(i)$
- S samples $\pi \leftarrow S_n$
- S permutes $\mathbf{x}$ by π to set OPRF inputs
- R samples a random PRF key k




Primary Construction

- Parties agree on a random vector $\mathbf{x}$ by computing $x_i = H(i)$
- S samples $\pi \leftarrow S_n$
- S permutes $\mathbf{x}$ by π to set OPRF inputs
- R samples a random PRF key k
- R sets $A_i = F_k(x_i)$ in plaintext




Primary Construction

- Parties agree on a random vector $\mathbf{x}$ by computing $x_i = H(i)$
- S samples $\pi \leftarrow S_n$
- S permutes $\mathbf{x}$ by π to set OPRF inputs
- R samples a random PRF key k
- R sets $A_i = F_k(x_i)$ in plaintext
- Parties set B and C as the OPRF outputs




Analysis of the Primary Construction




Analysis of the Primary Construction

- n calls to the OPRF




Analysis of the Primary Construction

- n calls to the OPRF
- $O(n\ell)$ total communication




Analysis of the Primary Construction

- n calls to the OPRF
- $O(n\ell)$ total communication
- Most efficient construction to date in practice



Analysis of the Primary Construction

- n calls to the OPRF
- $O(n\ell)$ total communication
- Most efficient construction to date in practice
- $\approx 4s$ per million elements



Sublinear Protocol

Sublinear Protocol

Treat the output of the primary construction as a vector of seeds to pseudorandom correlation generators.

Sublinear Protocol

Treat the output of the primary construction as a vector of seeds to pseudorandom correlation generators.



Sublinear Protocol

Treat the output of the primary construction as a vector of seeds to pseudorandom correlation generators.



- Let t be some sparsity parameter

Sublinear Protocol

Treat the output of the primary construction as a vector of seeds to pseudorandom correlation generators.



- Let t be some sparsity parameter
- Each $\log(2\ell/t)$ bits encodes a unit vector of length $2\ell/t$

Sublinear Protocol

Treat the output of the primary construction as a vector of seeds to pseudorandom correlation generators.



- Let t be some sparsity parameter
- Each $\log(2\ell/t)$ bits encodes a unit vector of length $2\ell/t$
- Compute each unit vector using a distributed point function (DPF)

Sublinear Protocol

Treat the output of the primary construction as a vector of seeds to pseudorandom correlation generators.



- Let t be some sparsity parameter
- Each $\log(2\ell/t)$ bits encodes a unit vector of length $2\ell/t$
- Compute each unit vector using a distributed point function (DPF)
- Concatenate t unit vectors to obtain a t -sparse vector of length 2ℓ

Sublinear Protocol

Treat the output of the primary construction as a vector of seeds to pseudorandom correlation generators.



- Let t be some sparsity parameter
- Each $\log(2\ell/t)$ bits encodes a unit vector of length $2\ell/t$
- Compute each unit vector using a distributed point function (DPF)
- Concatenate t unit vectors to obtain a t -sparse vector of length 2ℓ
- Compress into a pseudorandom vector of length ℓ

Sublinear Protocol

Treat the output of the primary construction as a vector of seeds to pseudorandom correlation generators.




- Let t be some sparsity parameter
- Each $\log(2\ell/t)$ bits encodes a unit vector of length $2\ell/t$
- Compute each unit vector using a distributed point function (DPF)
- Concatenate t unit vectors to obtain a t -sparse vector of length 2ℓ
- Compress into a pseudorandom vector of length ℓ
 - Reduction to syndrome decoding assumption

Sublinear Protocol

Treat the output of the primary construction as a vector of seeds to pseudorandom correlation generators.



- Let t be some sparsity parameter
- Each $\log(2\ell/t)$ bits encodes a unit vector of length $2\ell/t$
- Compute each unit vector using a distributed point function (DPF)
- Concatenate t unit vectors to obtain a t -sparse vector of length 2ℓ
- Compress into a pseudorandom vector of length ℓ
 - Reduction to syndrome decoding assumption



Analysis of the Sublinear Protocol

Analysis of the Sublinear Protocol

- $O(n \log \ell)$ total communication

Analysis of the Sublinear Protocol

- $O(n \log \ell)$ total communication
- Efficient only for very long ℓ

Analysis of the Sublinear Protocol

- $O(n \log \ell)$ total communication
- Efficient only for very long ℓ
- Crossover point is $\ell \approx 4000$

Analysis of the Sublinear Protocol

- $O(n \log \ell)$ total communication
- Efficient only for very long ℓ
- Crossover point is $\ell \approx 4000$
- Application: graph algorithms

Open Problems

Open Problems

Goal: Extend the line of work on PCGs into the domain of oblivious shuffling.

Open Problems

Goal: Extend the line of work on PCGs into the domain of oblivious shuffling.

We want to do to oblivious shuffling what pseudorandom correlation generators did to oblivious transfer.

Open Problems

Goal: Extend the line of work on PCGs into the domain of oblivious shuffling.

We want to do to oblivious shuffling what pseudorandom correlation generators did to oblivious transfer.

Concrete technical problems:

Open Problems

Goal: Extend the line of work on PCGs into the domain of oblivious shuffling.

We want to do to oblivious shuffling what pseudorandom correlation generators did to oblivious transfer.

Concrete technical problems:

- Sublinearity in ℓ (in practice)

Open Problems

Goal: Extend the line of work on PCGs into the domain of oblivious shuffling.

We want to do to oblivious shuffling what pseudorandom correlation generators did to oblivious transfer.

Concrete technical problems:

- Sublinearity in ℓ (in practice)
- Sublinearity in p

Open Problems

Goal: Extend the line of work on PCGs into the domain of oblivious shuffling.

We want to do to oblivious shuffling what pseudorandom correlation generators did to oblivious transfer.

Concrete technical problems:

- Sublinearity in ℓ (in practice)
- Sublinearity in p
- Sublinearity in n

Sublinearity in ℓ

Sublinearity in ℓ

Protocols exist for generating permutation correlations with sublinear communication in ℓ .

Sublinearity in ℓ

Protocols exist for generating permutation correlations with sublinear communication in ℓ .

- No implementation of the Pecený et al. construction (until now!)

Sublinearity in ℓ

Protocols exist for generating permutation correlations with sublinear communication in ℓ .

- No implementation of the Pecený et al. construction (until now!)
- Improve parameters to get practical efficiency

Sublinearity in ℓ

Protocols exist for generating permutation correlations with sublinear communication in ℓ .

- No implementation of the Pecený et al. construction (until now!)
- Improve parameters to get practical efficiency

Ongoing work: graph algorithms

Sublinearity in ℓ

Protocols exist for generating permutation correlations with sublinear communication in ℓ .

- No implementation of the Pecený et al. construction (until now!)
- Improve parameters to get practical efficiency

Ongoing work: graph algorithms

- Requires thousands of bits of permutation correlations

Sublinearity in ℓ

Protocols exist for generating permutation correlations with sublinear communication in ℓ .

- No implementation of the Pecený et al. construction (until now!)
- Improve parameters to get practical efficiency


Ongoing work: graph algorithms

- Requires thousands of bits of permutation correlations
- Existing constructions can be practical

Sublinearity in p

Sublinearity in p


Can we "chop up" a correlation with large ℓ into many correlations with small ℓ ?



Sublinearity in p

Can we "chop up" a correlation with large ℓ into many correlations with small ℓ ?


- Very challenging to achieve independence



Sublinearity in p

Can we "chop up" a correlation with large ℓ into many correlations with small ℓ ?

- Very challenging to achieve independence
- Can define an ideal functionality that captures the intended behavior




Sublinearity in p

Can we "chop up" a correlation with large ℓ into many correlations with small ℓ ?

- Very challenging to achieve independence
- Can define an ideal functionality that captures the intended behavior

Application: relational analytics



Sublinearity in n

Sublinearity in n

The holy grail of this research direction.

Sublinearity in n

The holy grail of this research direction.

- Make shuffling an efficient building block for *any* workload

Sublinearity in n

The holy grail of this research direction.

- Make shuffling an efficient building block for *any* workload

Preliminary ideas:

Sublinearity in n

The holy grail of this research direction.

- Make shuffling an efficient building block for *any* workload

Preliminary ideas:

- Generate many small correlations, then "stitch them together"

Sublinearity in n

The holy grail of this research direction.

- Make shuffling an efficient building block for *any* workload

Preliminary ideas:

- Generate many small correlations, then "stitch them together"
- Construct from a protocol with sublinearity in p

Sublinearity in n

The holy grail of this research direction.

- Make shuffling an efficient building block for *any* workload

Preliminary ideas:

- Generate many small correlations, then "stitch them together"
- Construct from a protocol with sublinearity in p
- Generate small correlations directly by materializing the permutation matrix

Scaling with Many Parties

Motivation

Motivation

Challenge: Suppose a secure computation involves thousands of parties.

Motivation

Challenge: Suppose a secure computation involves thousands of parties.

N Party Computation

Motivation

Challenge: Suppose a secure computation involves thousands of parties.

N Party Computation

- Large body of theoretical work

Motivation

Challenge: Suppose a secure computation involves thousands of parties.

N Party Computation

- Large body of theoretical work
- Prohibitive practical challenges


Motivation

Challenge: Suppose a secure computation involves thousands of parties.

N Party Computation

- Large body of theoretical work
- Prohibitive practical challenges

The Outsourced Setting



Motivation


Challenge: Suppose a secure computation involves thousands of parties.

N Party Computation

- Large body of theoretical work
- Prohibitive practical challenges

The Outsourced Setting

- Practical and efficient



Motivation


Challenge: Suppose a secure computation involves thousands of parties.

N Party Computation

- Large body of theoretical work
- Prohibitive practical challenges

The Outsourced Setting

- Practical and efficient
- Requires stronger trust assumptions



Motivation

Challenge: Suppose a secure computation involves thousands of parties.


N Party Computation

- Large body of theoretical work
- Prohibitive practical challenges

Goal: Scale *practical* MPC to thousands of parties.

The Outsourced Setting

- Practical and efficient
- Requires stronger trust assumptions



Secure Aggregation

Secure Aggregation

Many parties, very simple computations (sum).

Secure Aggregation

Many parties, very simple computations (sum).

- Ephemeral clients send encryptions to a powerful untrusted central server

Secure Aggregation

Many parties, very simple computations (sum).

- Ephemeral clients send encryptions to a powerful untrusted central server
- Design goal is non-interactivity

Secure Aggregation


Many parties, very simple computations (sum).

- Ephemeral clients send encryptions to a powerful untrusted central server
- Design goal is non-interactivity
- Naive approach is threshold homomorphic encryption

Secure Aggregation

Many parties, very simple computations (sum).

- Ephemeral clients send encryptions to a powerful untrusted central server
- Design goal is non-interactivity
- Naive approach is threshold homomorphic encryption
 - Vulnerable to "residual attack"




Secure Aggregation

Many parties, very simple computations (sum).

- Ephemeral clients send encryptions to a powerful untrusted central server
- Design goal is non-interactivity
- Naive approach is threshold homomorphic encryption
 - Vulnerable to "residual attack"

The emerging paradigm uses a small committee to atomically decrypt (OPA, Willow, TACITA).



A Star Topology

A Star Topology

- Damgård-Nielsen, 2007

Scalable and Unconditionally Secure Multiparty Computation

Ivan Damgård and Jesper Buus Nielsen*

Dept. of Computer Science, BRICS, Aarhus University

Abstract. We present a multiparty computation protocol that is unconditionally secure against adaptive and active adversaries, with communication complexity $\mathcal{O}(\mathcal{C}n)k + \mathcal{O}(Dn^2)k + \text{poly}(n\kappa)$, where $\mathcal{C}$ is the number of gates in the circuit, n is the number of parties, k is the bit-length of the elements of the field over which the computation is carried out, D is the multiplicative depth of the circuit, and κ is the security

A Star Topology


- Damgård-Nielsen, 2007
- Central untrusted "king" party

Scalable and Unconditionally Secure Multiparty Computation

Ivan Damgård and Jesper Buus Nielsen*

Dept. of Computer Science, BRICS, Aarhus University

Abstract. We present a multiparty computation protocol that is unconditionally secure against adaptive and active adversaries, with communication complexity $\mathcal{O}(\mathcal{C}n)k + \mathcal{O}(Dn^2)k + \text{poly}(n\kappa)$, where $\mathcal{C}$ is the number of gates in the circuit, n is the number of parties, k is the bit-length of the elements of the field over which the computation is carried out, D is the multiplicative depth of the circuit, and κ is the security



A Star Topology


- Damgård-Nielsen, 2007
- Central untrusted "king" party
- $O(n)$ communication per gate

Scalable and Unconditionally Secure Multiparty Computation

Ivan Damgård and Jesper Buus Nielsen*

Dept. of Computer Science, BRICS, Aarhus University

Abstract. We present a multiparty computation protocol that is unconditionally secure against adaptive and active adversaries, with communication complexity $\mathcal{O}(\mathcal{C}n)k + \mathcal{O}(Dn^2)k + \text{poly}(n\kappa)$, where $\mathcal{C}$ is the number of gates in the circuit, n is the number of parties, k is the bit-length of the elements of the field over which the computation is carried out, D is the multiplicative depth of the circuit, and κ is the security



A Star Topology


- Damgård-Nielsen, 2007
- Central untrusted "king" party
- $O(n)$ communication per gate
- Based on Shamir secret sharing

Scalable and Unconditionally Secure Multiparty Computation

Ivan Damgård and Jesper Buus Nielsen*

Dept. of Computer Science, BRICS, Aarhus University

Abstract. We present a multiparty computation protocol that is unconditionally secure against adaptive and active adversaries, with communication complexity $\mathcal{O}(\mathcal{C}n)k + \mathcal{O}(Dn^2)k + \text{poly}(n\kappa)$, where $\mathcal{C}$ is the number of gates in the circuit, n is the number of parties, k is the bit-length of the elements of the field over which the computation is carried out, D is the multiplicative depth of the circuit, and κ is the security



A Star Topology


- Damgård-Nielsen, 2007
- Central untrusted "king" party
- $O(n)$ communication per gate
- Based on Shamir secret sharing
- Requires an honest majority

Scalable and Unconditionally Secure Multiparty Computation

Ivan Damgård and Jesper Buus Nielsen*

Dept. of Computer Science, BRICS, Aarhus University

Abstract. We present a multiparty computation protocol that is unconditionally secure against adaptive and active adversaries, with communication complexity $\mathcal{O}(\mathcal{C}n)k + \mathcal{O}(Dn^2)k + \text{poly}(n\kappa)$, where $\mathcal{C}$ is the number of gates in the circuit, n is the number of parties, k is the bit-length of the elements of the field over which the computation is carried out, D is the multiplicative depth of the circuit, and κ is the security



DN07 Multiplication

DN07 Multiplication

- $[x]$ denotes a degree- t Shamir sharing

DN07 Multiplication

- $[x]$ denotes a degree- t Shamir sharing
- $[[x]]$ denotes a degree- $2t$ Shamir sharing

DN07 Multiplication

- $[x]$ denotes a degree- t Shamir sharing
- $[[x]]$ denotes a degree- $2t$ Shamir sharing
- Local multiplication converts t sharing to $2t$ sharing

DN07 Multiplication


- $[x]$ denotes a degree- t Shamir sharing
- $[[x]]$ denotes a degree- $2t$ Shamir sharing
- Local multiplication converts t sharing to $2t$ sharing
- We need to perform a degree reduction

DN07 Multiplication

- $[x]$ denotes a degree- t Shamir sharing
- $[[x]]$ denotes a degree- $2t$ Shamir sharing
- Local multiplication converts t sharing to $2t$ sharing
- We need to perform a degree reduction
- Multiplication consumes a double sharing $([r], [[r]])$


DN07 Multiplication

- $[x]$ denotes a degree- t Shamir sharing
- $[[x]]$ denotes a degree- $2t$ Shamir sharing
- Local multiplication converts t sharing to $2t$ sharing
- We need to perform a degree reduction
- Multiplication consumes a double sharing $([r], [[r]])$




DN07 Multiplication

- $[x]$ denotes a degree- t Shamir sharing
- $[[x]]$ denotes a degree- $2t$ Shamir sharing
- Local multiplication converts t sharing to $2t$ sharing
- We need to perform a degree reduction
- Multiplication consumes a double sharing $([r], [[r]])$



DN07 Multiplication


- $[x]$ denotes a degree- t Shamir sharing
- $[[x]]$ denotes a degree- $2t$ Shamir sharing
- Local multiplication converts t sharing to $2t$ sharing
- We need to perform a degree reduction
- Multiplication consumes a double sharing $([r], [[r]])$



$\text{Reconstruct}([S_1], \dots, [S_n]) \rightarrow \text{secret} + \text{mask}$

DN07 Multiplication


- $[x]$ denotes a degree- t Shamir sharing
- $[[x]]$ denotes a degree- $2t$ Shamir sharing
- Local multiplication converts t sharing to $2t$ sharing
- We need to perform a degree reduction
- Multiplication consumes a double sharing $([r], [[r]])$



$\text{reshare}(\text{secret} + \text{mask}) \rightarrow [S_1], \dots, [S_n]$

DN07 Multiplication

- $[x]$ denotes a degree- t Shamir sharing
- $[[x]]$ denotes a degree- $2t$ Shamir sharing
- Local multiplication converts t sharing to $2t$ sharing
- We need to perform a degree reduction
- Multiplication consumes a double sharing $([r], [[r]])$



Extending to More Complex Functions

Extending to More Complex Functions

Liu et al. (Usenix Security 2024)

Scalable Multi-Party Computation Protocols for Machine Learning in the Honest-Majority Setting

Fengrun Liu
*University of Science and Technology of China &
Shanghai Qi Zhi Institute*

Xiang Xie
Shanghai Qi Zhi Institute & PADO Labs

Yu Yu
*Shanghai Jiao Tong University &
State Key Laboratory of Cryptology, P. O. Box 5159, Beijing, 100878, China*

Abstract

In this paper, we present a novel and scalable multi-party computation (MPC) protocol tailored for privacy-preserving machine learning (PPML) with semi-honest security in the honest-majority setting. Our protocol utilizes the Damgård-Nielsen (Crypto'07) protocol with Mersenne prime fields.

Secure Multi-Party Computation (MPC) [3, 21, 47] enables a group of n parties to collectively compute a function on their private inputs while preserving the privacy of the inputs. MPC offers a promising approach to address these privacy challenges in ML. It enables privacy-preserving machine learning (PPML) by performing secure computations on dis-

Extending to More Complex Functions

Liu et al. (Usenix Security 2024)

- Based on the DN07 protocol

Scalable Multi-Party Computation Protocols for Machine Learning in the Honest-Majority Setting

Fengrun Liu
*University of Science and Technology of China &
Shanghai Qi Zhi Institute*

Xiang Xie
Shanghai Qi Zhi Institute & PADO Labs

Yu Yu
*Shanghai Jiao Tong University &
State Key Laboratory of Cryptology, P. O. Box 5159, Beijing, 100878, China*

Abstract

In this paper, we present a novel and scalable multi-party computation (MPC) protocol tailored for privacy-preserving machine learning (PPML) with semi-honest security in the honest-majority setting. Our protocol utilizes the Damgård-Nielsen (DN07) protocol with Mersenne prime fields.

Secure Multi-Party Computation (MPC) [3, 21, 47] enables a group of n parties to collectively compute a function on their private inputs while preserving the privacy of the inputs. MPC offers a promising approach to address these privacy challenges in ML. It enables privacy-preserving machine learning (PPML) by performing secure computations on dis-

Extending to More Complex Functions

Liu et al. (Usenix Security 2024)

- Based on the DN07 protocol
- Uses the properties of Mersenne prime fields

Scalable Multi-Party Computation Protocols for Machine Learning in the Honest-Majority Setting

Fengrun Liu
University of Science and Technology of China & Shanghai Qi Zhi Institute

Xiang Xie
Shanghai Qi Zhi Institute & PADO Labs

Yu Yu
Shanghai Jiao Tong University & State Key Laboratory of Cryptology, P. O. Box 5159, Beijing, 100878, China

Abstract

In this paper, we present a novel and scalable multi-party computation (MPC) protocol tailored for privacy-preserving machine learning (PPML) with semi-honest security in the honest-majority setting. Our protocol utilizes the Damgård-Nielsen (DN07) protocol with Mersenne prime fields.

Secure Multi-Party Computation (MPC) [3, 21, 47] enables a group of n parties to collectively compute a function on their private inputs while preserving the privacy of the inputs. MPC offers a promising approach to address these privacy challenges in ML. It enables privacy-preserving machine learning (PPML) by performing secure computations on dis-

Extending to More Complex Functions

Liu et al. (Usenix Security 2024)

- Based on the DN07 protocol
- Uses the properties of Mersenne prime fields
- Bitwise operations on arithmetic secret shares

Scalable Multi-Party Computation Protocols for Machine Learning in the Honest-Majority Setting

Fengrun Liu
*University of Science and Technology of China &
Shanghai Qi Zhi Institute*

Xiang Xie
Shanghai Qi Zhi Institute & PADO Labs

Yu Yu
*Shanghai Jiao Tong University &
State Key Laboratory of Cryptology, P. O. Box 5159, Beijing, 100878, China*

Abstract

In this paper, we present a novel and scalable multi-party computation (MPC) protocol tailored for privacy-preserving machine learning (PPML) with semi-honest security in the honest-majority setting. Our protocol utilizes the Damgård-Nielsen (DN07) protocol with Mersenne prime fields.

Secure Multi-Party Computation (MPC) [3, 21, 47] enables a group of n parties to collectively compute a function on their private inputs while preserving the privacy of the inputs. MPC offers a promising approach to address these privacy challenges in ML. It enables privacy-preserving machine learning (PPML) by performing secure computations on dis-

Extending to More Complex Functions

Liu et al. (Usenix Security 2024)

- Based on the DN07 protocol
- Uses the properties of Mersenne prime fields
- Bitwise operations on arithmetic secret shares
- Targets machine learning inference

Scalable Multi-Party Computation Protocols for Machine Learning in the Honest-Majority Setting

Fengrun Liu
University of Science and Technology of China & Shanghai Qi Zhi Institute

Xiang Xie
Shanghai Qi Zhi Institute & PADO Labs

Yu Yu
Shanghai Jiao Tong University & State Key Laboratory of Cryptology, P. O. Box 5159, Beijing, 100878, China

Abstract

In this paper, we present a novel and scalable multi-party computation (MPC) protocol tailored for privacy-preserving machine learning (PPML) with semi-honest security in the honest-majority setting. Our protocol utilizes the Damgård-Nielsen (DN07) protocol with Mersenne prime fields.

Secure Multi-Party Computation (MPC) [3, 21, 47] enables a group of n parties to collectively compute a function on their private inputs while preserving the privacy of the inputs. MPC offers a promising approach to address these privacy challenges in ML. It enables privacy-preserving machine learning (PPML) by performing secure computations on dis-

Extending to More Complex Functions

Liu et al. (Usenix Security 2024)

- Based on the DN07 protocol
- Uses the properties of Mersenne prime fields
- Bitwise operations on arithmetic secret shares
- Targets machine learning inference
- Optimizes for round complexity

Scalable Multi-Party Computation Protocols for Machine Learning in the Honest-Majority Setting

Fengrun Liu
*University of Science and Technology of China &
Shanghai Qi Zhi Institute*

Xiang Xie
Shanghai Qi Zhi Institute & PADO Labs

Yu Yu
*Shanghai Jiao Tong University &
State Key Laboratory of Cryptology, P. O. Box 5159, Beijing, 100878, China*

Abstract

In this paper, we present a novel and scalable multi-party computation (MPC) protocol tailored for privacy-preserving machine learning (PPML) with semi-honest security in the honest-majority setting. Our protocol utilizes the Damgård-Nielsen (DN07) protocol with Mersenne prime fields.

Secure Multi-Party Computation (MPC) [3, 21, 47] enables a group of n parties to collectively compute a function on their private inputs while preserving the privacy of the inputs. MPC offers a promising approach to address these privacy challenges in ML. It enables privacy-preserving machine learning (PPML) by performing secure computations on dis-

Extending to More Complex Functions

Liu et al. (Usenix Security 2024)

- Based on the DN07 protocol
- Uses the properties of Mersenne prime fields
- Bitwise operations on arithmetic secret shares
- Targets machine learning inference
- Optimizes for round complexity
- Shows results with up to 63 parties

Scalable Multi-Party Computation Protocols for Machine Learning in the Honest-Majority Setting

Fengrun Liu
*University of Science and Technology of China &
Shanghai Qi Zhi Institute*

Xiang Xie
Shanghai Qi Zhi Institute & PADO Labs

Yu Yu
*Shanghai Jiao Tong University &
State Key Laboratory of Cryptology, P. O. Box 5159, Beijing, 100878, China*

Abstract

In this paper, we present a novel and scalable multi-party computation (MPC) protocol tailored for privacy-preserving machine learning (PPML) with semi-honest security in the honest-majority setting. Our protocol utilizes the Damgård-Nielsen (DN07) protocol with Mersenne prime fields.

Secure Multi-Party Computation (MPC) [3, 21, 47] enables a group of n parties to collectively compute a function on their private inputs while preserving the privacy of the inputs. MPC offers a promising approach to address these privacy challenges in ML. It enables privacy-preserving machine learning (PPML) by performing secure computations on dis-

Liu et al. Truncation Protocol

Liu et al. Truncation Protocol

- Division by a public power of two

Liu et al. Truncation Protocol

- Division by a public power of two

Almost all truncation protocols either

Liu et al. Truncation Protocol

- Division by a public power of two

Almost all truncation protocols either

- Require an arithmetic-to-binary conversion

Liu et al. Truncation Protocol

- Division by a public power of two

Almost all truncation protocols either

- Require an arithmetic-to-binary conversion
- Require "slack bits" to prevent a "catastrophic" error

Liu et al. Truncation Protocol

- Division by a public power of two

Almost all truncation protocols either

- Require an arithmetic-to-binary conversion
- Require "slack bits" to prevent a "catastrophic" error

The Liu et al. protocol is both constant-round and requires only 1 bit of slack.

Liu et al. Truncation Protocol

- Division by a public power of two

Almost all truncation protocols either

- Require an arithmetic-to-binary conversion
- Require "slack bits" to prevent a "catastrophic" error

The Liu et al. protocol is both constant-round and requires only 1 bit of slack.

- In a Mersenne prime field, the catastrophic error is predictable and can be corrected for

Liu et al. Truncation Protocol

- Division by a public power of two

Almost all truncation protocols either

- Require an arithmetic-to-binary conversion
- Require "slack bits" to prevent a "catastrophic" error

The Liu et al. protocol is both constant-round and requires only 1 bit of slack.

- In a Mersenne prime field, the catastrophic error is predictable and can be corrected for

The paper simultaneously minimizes total work and makes base operations constant-round.

Liu et al. in Practice

Liu et al. in Practice

LAN Setting

Liu et al. in Practice

LAN Setting

- Bottlenecked by local computation on the king party

Liu et al. in Practice

LAN Setting

- Bottlenecked by local computation on the king party
- King needs to perform polynomial interpolation for each gate

Liu et al. in Practice

LAN Setting

- Bottlenecked by local computation on the king party
- King needs to perform polynomial interpolation for each gate

Liu et al. in Practice

LAN Setting

- Bottlenecked by local computation on the king party
- King needs to perform polynomial interpolation for each gate

WAN Setting

Liu et al. in Practice

LAN Setting

- Bottlenecked by local computation on the king party
- King needs to perform polynomial interpolation for each gate

WAN Setting

- Latency constrained

Liu et al. in Practice

LAN Setting

- Bottlenecked by local computation on the king party
- King needs to perform polynomial interpolation for each gate

WAN Setting

- Latency constrained
- Uses 5% of available bandwidth

Scaling the King Party

Scaling the King Party

Proposal: Lean into the king as the bottleneck.

Scaling the King Party

Proposal: Lean into the king as the bottleneck.

- Powerful central server, weak ephemeral clients

Scaling the King Party

Proposal: Lean into the king as the bottleneck.

- Powerful central server, weak ephemeral clients
- One logical king, multiple servers



Scaling the King Party

Proposal: Lean into the king as the bottleneck.

- Powerful central server, weak ephemeral clients
- One logical king, multiple servers
- GPU acceleration



Tolerating Dropouts


Tolerating Dropouts

Problem: Slowest connection becomes the bottleneck.

Tolerating Dropouts

Problem: Slowest connection becomes the bottleneck.

Proposal: Embrace the threshold scheme to tolerate dropouts.




Tolerating Dropouts

Problem: Slowest connection becomes the bottleneck.

Proposal: Embrace the threshold scheme to tolerate dropouts.

- Subset of Fluid MPC




Tolerating Dropouts

Problem: Slowest connection becomes the bottleneck.

Proposal: Embrace the threshold scheme to tolerate dropouts.

- Subset of Fluid MPC
- Semi-honest security is straightforward




Tolerating Dropouts

Problem: Slowest connection becomes the bottleneck.

Proposal: Embrace the threshold scheme to tolerate dropouts.

- Subset of Fluid MPC
- Semi-honest security is straightforward
- Malicious security is more complex



Summary

Summary

Vision: Efficient secure computation involving millions of parties for complex functions.

Summary

Vision: Efficient secure computation involving millions of parties for complex functions.

- Oblivious shuffle as a way to increase the expressivity of MPC

Summary

Vision: Efficient secure computation involving millions of parties for complex functions.

- Oblivious shuffle as a way to increase the expressivity of MPC
 - Incorporating pseudorandom correlation generators

Summary

Vision: Efficient secure computation involving millions of parties for complex functions.

- Oblivious shuffle as a way to increase the expressivity of MPC
 - Incorporating pseudorandom correlation generators
- Approaches to scale moderately complex workloads to thousands of parties

Summary

Vision: Efficient secure computation involving millions of parties for complex functions.

- Oblivious shuffle as a way to increase the expressivity of MPC
 - Incorporating pseudorandom correlation generators
- Approaches to scale moderately complex workloads to thousands of parties
 - Embracing the star topology

Summary

Vision: Efficient secure computation involving millions of parties for complex functions.

- Oblivious shuffle as a way to increase the expressivity of MPC
 - Incorporating pseudorandom correlation generators
- Approaches to scale moderately complex workloads to thousands of parties
 - Embracing the star topology
 - Tolerating dropouts

Thank You!